

Yield Prediction and Analysis

About Dataset

This dataset encompasses agricultural data for multiple crops cultivated across various states in India from the year 2011 till 2020. The dataset provides crucial features related to crop yield prediction, including crop types, crop years, cropping seasons, states, areas under cultivation, production quantities, annual rainfall, fertilizer usage, pesticide usage, and calculated yields.

Data Fields

- Crop: The name of the crop cultivated.
- Crop_Year: The year in which the crop was grown.
- Season: The specific cropping season (e.g., Kharif, Rabi, Whole Year).
- State: The Indian state where the crop was cultivated.
- Area: The total land area (in hectares) under cultivation for the specific crop.
- Production: The quantity of crop production (in metric tons).
- Annual_Rainfall: The annual rainfall received in the crop-growing region (in mm).
- Fertilizer: The total amount of fertilizer used for the crop (in kilograms).
- Pesticide: The total amount of pesticide used for the crop (in kilograms).
- Yield: The calculated crop yield (production per unit area).

Importing Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

Importing Agriculture Crop Yield Dataset and Initial Observations

```
data = pd.read_csv('crop_yield.csv')
data.head()
```

\	Crop	Crop_Year	Season	State	Area	Production
0	Arecanut	2011	Kharif	Puducherry	60.0	77
1	Bajra	2011	Kharif	Puducherry	11.0	29
2	Bajra	2011	Summer	Puducherry	20.0	51
3	Banana	2011	Kharif	Puducherry	266.0	3263
4	Black pepper	2011	Kharif	Puducherry	11.0	5

	Annual_Rainfall	Fertilizer	Pesticide	Yield
0	1434.5875	10051.20	19.80	1.280000
1	1434.5875	1842.72	3.63	2.640000
2	1434.5875	3350.40	6.60	2.550000
3	1434.5875	44560.32	87.78	10.903333
4	1434.5875	1842.72	3.63	0.450000

```
data.tail()
```

	Crop	Crop_Year	Season	State	Area	
Production \						
9017	Sugarcane	2018	Winter	Odisha	6778.0	417672
9018	Urad	2018	Autumn	Odisha	13720.0	3583
9019	Urad	2018	Summer	Odisha	4571.0	2336
9020	Urad	2018	Winter	Odisha	39560.0	13123
9021	Wheat	2018	Summer	Odisha	147.0	268

	Annual_Rainfall	Fertilizer	Pesticide	Yield
9017	1635.9	1099391.6	2372.30	57.584545
9018	1635.9	2225384.0	4802.00	0.336667
9019	1635.9	741416.2	1599.85	0.469091
9020	1635.9	6416632.0	13846.00	0.352759
9021	1635.9	23843.4	51.45	1.825000

```
data.shape
```

```
(9022, 10)
```

We understand that our data has 9022 rows and 10 features to work with

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 9022 entries, 0 to 9021
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Crop                  9022 non-null   object
1   Crop_Year             9022 non-null   int64
2   Season                9022 non-null   object
3   State                 9022 non-null   object
4   Area                  9022 non-null   float64
5   Production            9022 non-null   int64
6   Annual_Rainfall       9022 non-null   float64
```

```

7   Fertilizer      9022 non-null   float64
8   Pesticide      9022 non-null   float64
9   Yield          9022 non-null   float64
dtypes: float64(5), int64(2), object(3)
memory usage: 705.0+ KB

```

This information tells us that, out of the 10 features in our dataset 3 are object type else all are numerical data

```
data.describe(include='all').T
```

	count	unique	top	freq	mean	\
Crop	9022	55	Rice	527	NaN	
Crop_Year	9022.0	NaN	NaN	NaN	2015.182997	
Season	9022	6	Kharif	3769	NaN	
State	9022	30	Karnataka	621	NaN	
Area	9022.0	NaN	NaN	NaN	159062.678774	
Production	9022.0	NaN	NaN	NaN	17051337.986588	
Annual_Rainfall	9022.0	NaN	NaN	NaN	1435.353493	
Fertilizer	9022.0	NaN	NaN	NaN	25052366.189199	
Pesticide	9022.0	NaN	NaN	NaN	53458.906243	
Yield	9022.0	NaN	NaN	NaN	82.457035	
	std	min	25%	50%	\	
Crop	NaN	NaN	NaN	NaN	NaN	
Crop_Year	2.582179	2011.0	2013.0	2015.0		
Season	NaN	NaN	NaN	NaN	NaN	
State	NaN	NaN	NaN	NaN	NaN	
Area	630033.776956	0.8	1036.25	7000.0		
Production	270796230.070602	0.0	1052.5	11117.5		
Annual_Rainfall	785.527676	301.3	956.2	1235.6		
Fertilizer	99436190.084488	120.768	163052.9676	1100093.175		
Pesticide	212682.696615	0.264	347.5925	2335.53		
Yield	923.014095	0.0	0.683333	1.1255		
	75%	max				
Crop	NaN	NaN				
Crop_Year	2017.0	2020.0				
Season	NaN	NaN				

State	NaN	NaN
Area	56567.5	10216517.0
Production	108095.75	6200900000.0
Annual_Rainfall	1593.9	5649.1
Fertilizer	8878633.35	1754788960.0
Pesticide	18940.05	3780111.29
Yield	2.699125	21105.0

```
data.isnull().sum()
```

```
Crop      0
Crop_Year 0
Season    0
State     0
Area      0
Production 0
Annual_Rainfall 0
Fertilizer 0
Pesticide 0
Yield     0
dtype: int64
```

This shows that our data has no null values in it

```
data.head()
```


	Crop	Crop_Year	Season	State	Area	Production
0	Arecanut	2011	Kharif	Puducherry	60.0	77
1	Bajra	2011	Kharif	Puducherry	11.0	29
2	Bajra	2011	Summer	Puducherry	20.0	51
3	Banana	2011	Kharif	Puducherry	266.0	3263
4	Black pepper	2011	Kharif	Puducherry	11.0	5

	Annual_Rainfall	Fertilizer	Pesticide	Yield
0	1434.5875	10051.20	19.80	1.280000
1	1434.5875	1842.72	3.63	2.640000
2	1434.5875	3350.40	6.60	2.550000
3	1434.5875	44560.32	87.78	10.903333
4	1434.5875	1842.72	3.63	0.450000

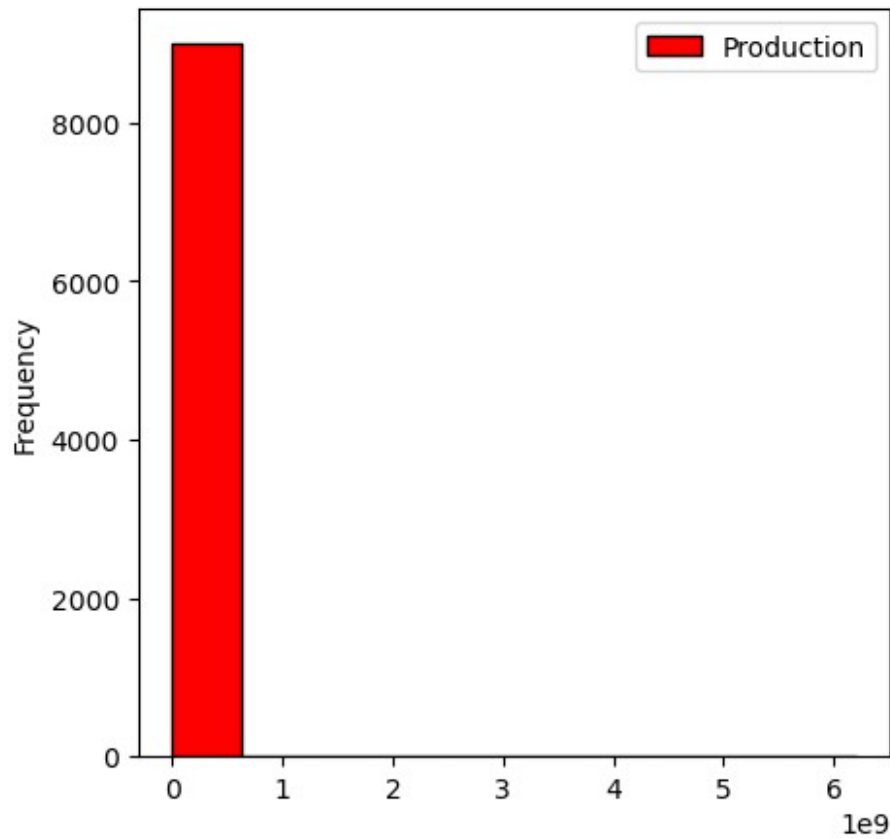
Exploratory Data Analysis

```
data['Area'].plot(kind='hist',figsize=(5,5),color='red',edgecolor='black')
```

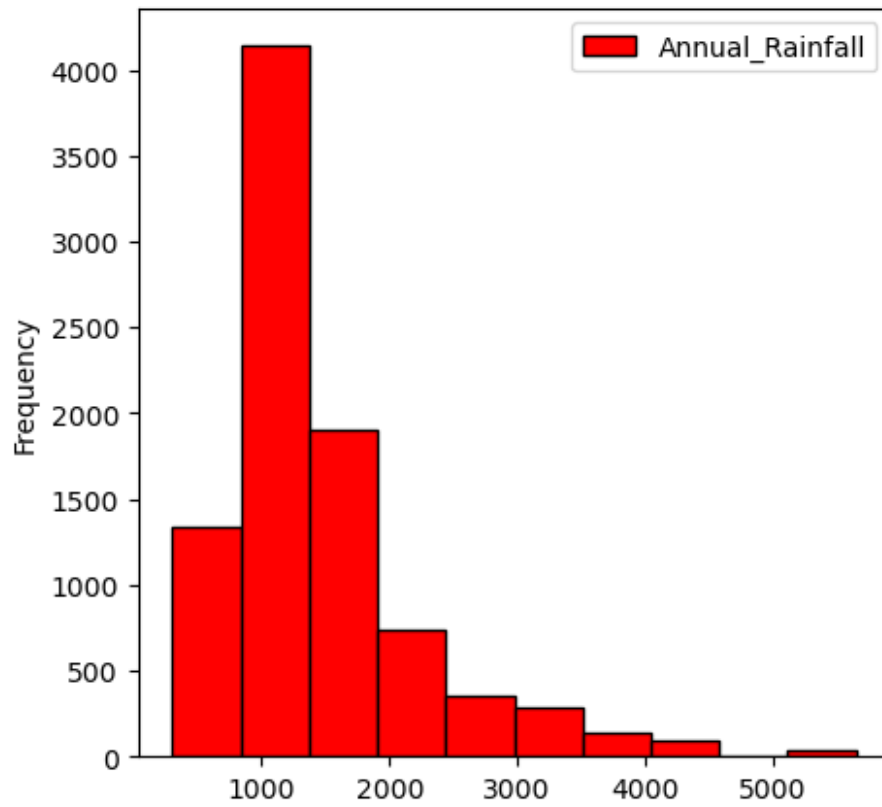
```
plt.legend()  
plt.show()
```



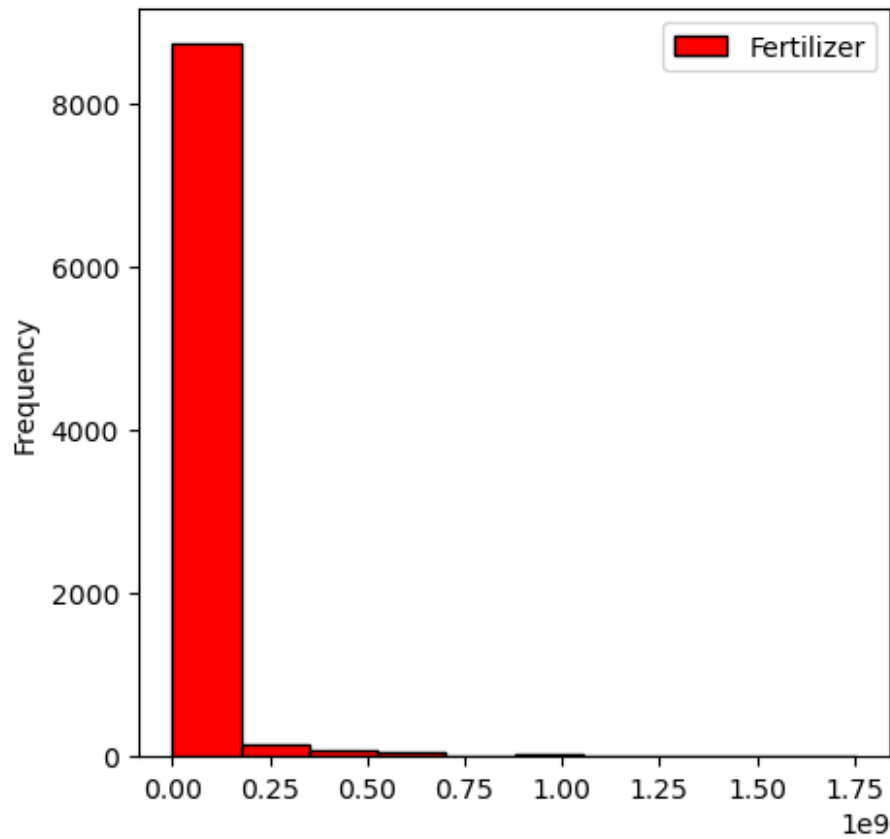
```
data['Production'].plot(kind='hist',figsize=(5,5),color='red',edgecolor='black')  
plt.legend()  
plt.show()
```



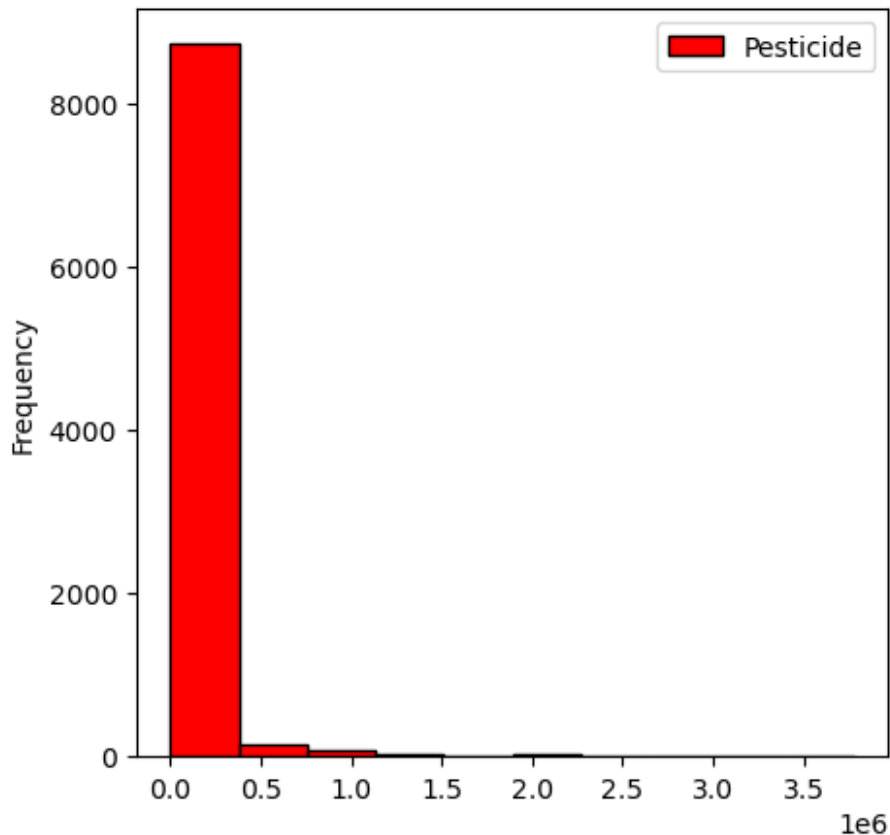
```
data['Annual_Rainfall'].plot(kind='hist',figsize=(5,5),color='red',edgecolor='black')  
plt.legend()  
plt.show()
```



```
data['Fertilizer'].plot(kind='hist',figsize=(5,5),color='red',edgecolor='black')  
plt.legend()  
plt.show()
```



```
data['Pesticide'].plot(kind='hist',figsize=(5,5),color='red',edgecolor='black')  
plt.legend()  
plt.show()
```

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.metrics import mean_squared_error, r2_score

# List of categorical and numerical columns
categorical_columns = ['Crop', 'Season', 'State']
numerical_columns = ['Crop_Year', 'Area', 'Annual_Rainfall',
                     'Fertilizer', 'Pesticide']

# Encoding categorical variables using one-hot encoding
data_encoded = pd.get_dummies(data, columns=categorical_columns,
                              drop_first=True)

# Standardizing numerical features
scaler = StandardScaler()
data_encoded[numerical_columns] =
scaler.fit_transform(data_encoded[numerical_columns])

# Normalizing the target variable (Yield)
target_scaler = MinMaxScaler()
data_encoded['Yield'] =
target_scaler.fit_transform(data_encoded[['Yield']])
```

```

# Splitting features and target variable
X = data_encoded.drop(columns=['Yield']) # Features
y = data_encoded['Yield'] # Target

# Train-test splitting
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Training the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Making predictions
y_pred = model.predict(X_test)

# Model eEvaluation
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

# Print evaluation metrics
print(f"R2 Score: {r2:.4f}")
print(f"Mean Squared Error (MSE): {mse:.4f}")
print(f"Root Mean Squared Error (RMSE): {rmse:.4f}")

R2 Score: 0.8541
Mean Squared Error (MSE): 0.0004
Root Mean Squared Error (RMSE): 0.0191

sorted_indices = np.argsort(y_test)
y_test_sorted = y_test.iloc[sorted_indices]
y_pred_sorted = y_pred[sorted_indices]

# Plotting the Actual vs. Predicted with Fit Line
plt.figure(figsize=(8, 6))
plt.plot(y_test_sorted, y_pred_sorted, color='blue', label='Fit Line')
plt.scatter(range(len(y_test_sorted)), y_test_sorted, color='red',
label='Actual Values', alpha=0.6)
plt.scatter(range(len(y_pred_sorted)), y_pred_sorted, color='green',
label='Predicted Values', alpha=0.6)
plt.xlabel("Index")
plt.ylabel("Yield (Normalized)")
plt.title("Fit Line for Linear Regression")
plt.legend()
plt.grid(True)
plt.show()

```

